

Week 10:

Natural Language Processing (NLP)

Vinay P. Namboodiri

Topic 8: Word Embedding

- GloVe

GloVe

- Proposed by Pennington et al. 2014.
- Unlike Word2vec, GloVe does not rely just on local statistics (local context information of words), but incorporates global statistics (word co-occurrence matrix) to obtain word vectors.



GloVe

Counts the co-occurrence between words w_k, w_i

$$P(w_k|w_i) = \frac{C(w_k, w_i)}{C(w_i)}$$

Probability and Ratio	$k = \text{solid}$	$k = \text{gas}$	$k = \text{water}$	$k = \text{fashion}$
$P(k \text{ice})$	1.9×10^{-4}	6.6×10^{-5}	3.0×10^{-3}	1.7×10^{-5}
$P(k \text{steam})$	2.2×10^{-5}	7.8×10^{-4}	2.2×10^{-3}	1.8×10^{-5}
$P(k \text{ice})/P(k \text{steam})$	8.9	8.5×10^{-2}	1.36	0.96

“solid” is more related to
“ice” than “steam”

$$F(w_i, w_j, w_k) = \frac{P(w_k|w_i)}{P(w_k|w_j)}$$

“Water” is equally (ir)relevant to
“ice” and “steam”

Intuition: co-occurrence probabilities ratios gathers more information than the raw probabilities and better capture relevant information about words' relationship

What is the function $F(\cdot)$?

$$F(w_i, w_j, w_k) = \frac{P(w_k|w_i)}{P(w_k|w_j)}$$

Since the goal is to learn meaningful word vectors, $F(\cdot)$ is designed to be a function of the linear difference between two words w_i and w_j

$$F((w_i - w_j)^T w_k) = \frac{P(w_k|w_i)}{P(w_k|w_j)}$$

The final solution is to $F(\cdot)$ as an **exponential** function.

$$F((w_i - w_j)^T w_k) = \exp((w_i - w_j)^T w_k) = \frac{\exp(w_i^T w_k)}{\exp(w_j^T w_k)} = \frac{P(w_k|w_i)}{P(w_k|w_j)}$$

Loss function of GloVe

Replace it with a bias term b_i

$$w_i^T w_k = \log \frac{C(w_k, w_i)}{C(w_i)} = \log C(w_k, w_i) - \log C(w_i)$$

$$\log C(w_k, w_i) = w_i^T w_k + b_i + b_k$$

To keep the symmetric form,
we also add in a bias term b_k

The loss function for the GloVe model is designed to preserve the above formula by minimizing the sum of the squared errors

$$\mathcal{L} = \sum_{i=1, j=1}^V (w_i^T w_j + b_i + b_j - \log C(w_j, w_i))^2$$

Add a weighting function $f()$ that is used to downweight the importance of very frequent co-occurrences

$$\mathcal{L} = \sum_{i=1, j=1}^V f(C(w_j, w_i)) (w_i^T w_j + b_i + b_j - \log C(w_j, w_i))^2$$

Loss function of GloVe

To penalize the difference between the dot product of two word vectors and the logarithm of the co-occurrence count, with the bias terms added.

$$\mathcal{L} = \sum_{i=1, j=1}^V f(C(w_j, w_i))(w_i^T w_j + b_i + b_j - \log C(w_j, w_i))^2$$

w_i, w_j are the word vectors for words i and j

b_i, b_j are bias terms

$$f(c) = \begin{cases} \left(\frac{c}{c_{\max}}\right)^\alpha & \text{if } c < c_{\max}, c_{\max} \text{ is adjustable.} \\ 1 & \text{if otherwise} \end{cases}$$

Nearest words to
frog:

1. frogs
2. toad
3. litoria
4. leptodactylidae
5. rana
6. lizard
7. eleutherodactylus



litoria



leptodactylidae



rana



eleutherodactylus

Advantages of GloVe

- Computationally efficient
- Scales well to large datasets
- Produces embedding that capture both syntactic and semantic relationships between words.

GloVe vs. Word2Vec

- Both are popular algorithms for generating word embeddings.
- Both are unsupervised learning algorithms
- Both are able to capture semantic relationships between words.

Word2Vec

Use a neural network to learn embedding
Focus more on local context
Is able to handle larger corpora of text

GloVe

Based on co-occurrence matrix
Capture global relationships between words.
Is generally faster than Word2Vec

Reference

- Pennington, J., Socher, R., & Manning, C. D. (2014). GloVe: Global Vectors for Word Representation. Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP).
- <https://nlp.stanford.edu/projects/glove/>
- Lectures of CS224n: NLP with Deep learning.
<https://web.stanford.edu/class/archive/cs/cs224n/cs224n.1184/syllabus.html>
- Matyas Amrouche's blog: Word embedding (Part II).
<https://towardsdatascience.com/word-embedding-part-ii-intuition-and-some-maths-to-understand-end-to-end-glove-model-9b08e6bf5c06>
- Lilian Weng's blog: Learning Word Embedding.
<https://lilianweng.github.io/posts/2017-10-15-word-embedding/>

Week 11: Domain Adaptation

Vinay P. Namboodiri

Introduction

- So far, most of the techniques we have considered assume the availability of full supervision for training through a training dataset
- However, in many practical scenarios, this is not true
- For instance, for an autonomous car, you may have trained a pedestrian detection algorithm in summer
- Then it snows.....

Problem



Figure : Typical Computer Vision Dataset

Generally, training and test instances are chosen from the same dataset and hence they are from same probability distribution. Also labels are free of noise, objects are centered and background clutter is less.

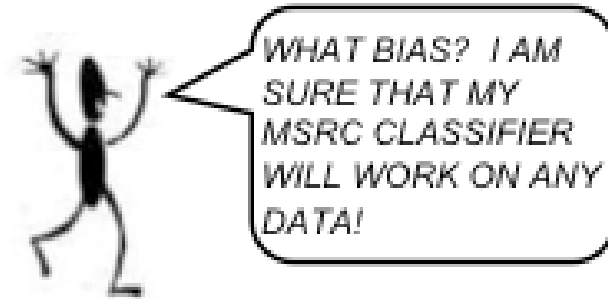
Challenge



Figure : Real World Computer Vision Dataset

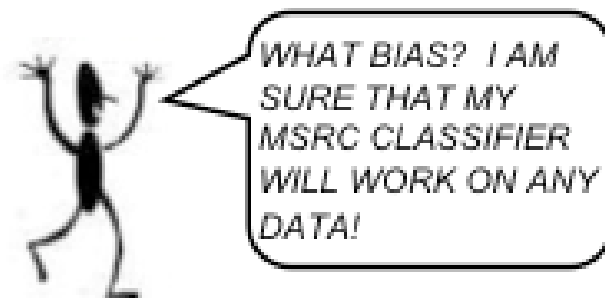
Real world data is noisy. Multiple instances of different object can be present in an image and also usually much more clutter is there.

Problem

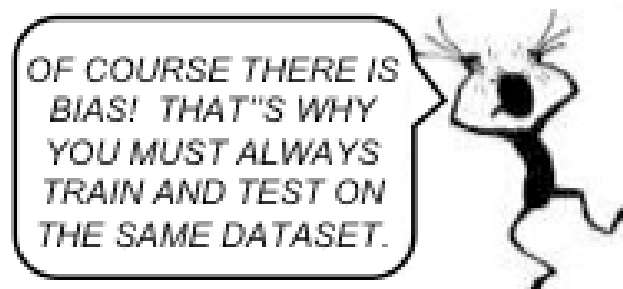


1. Denial

Problem

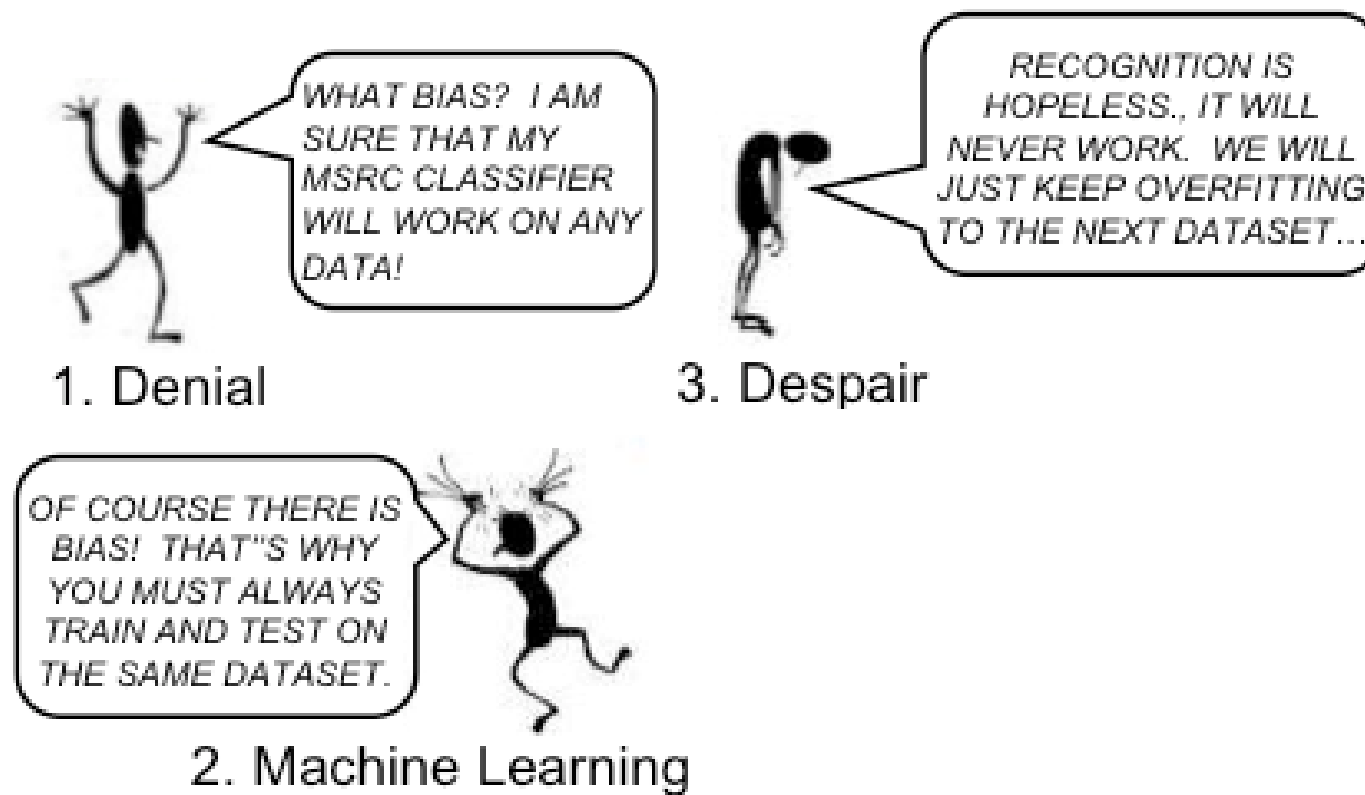


1. Denial

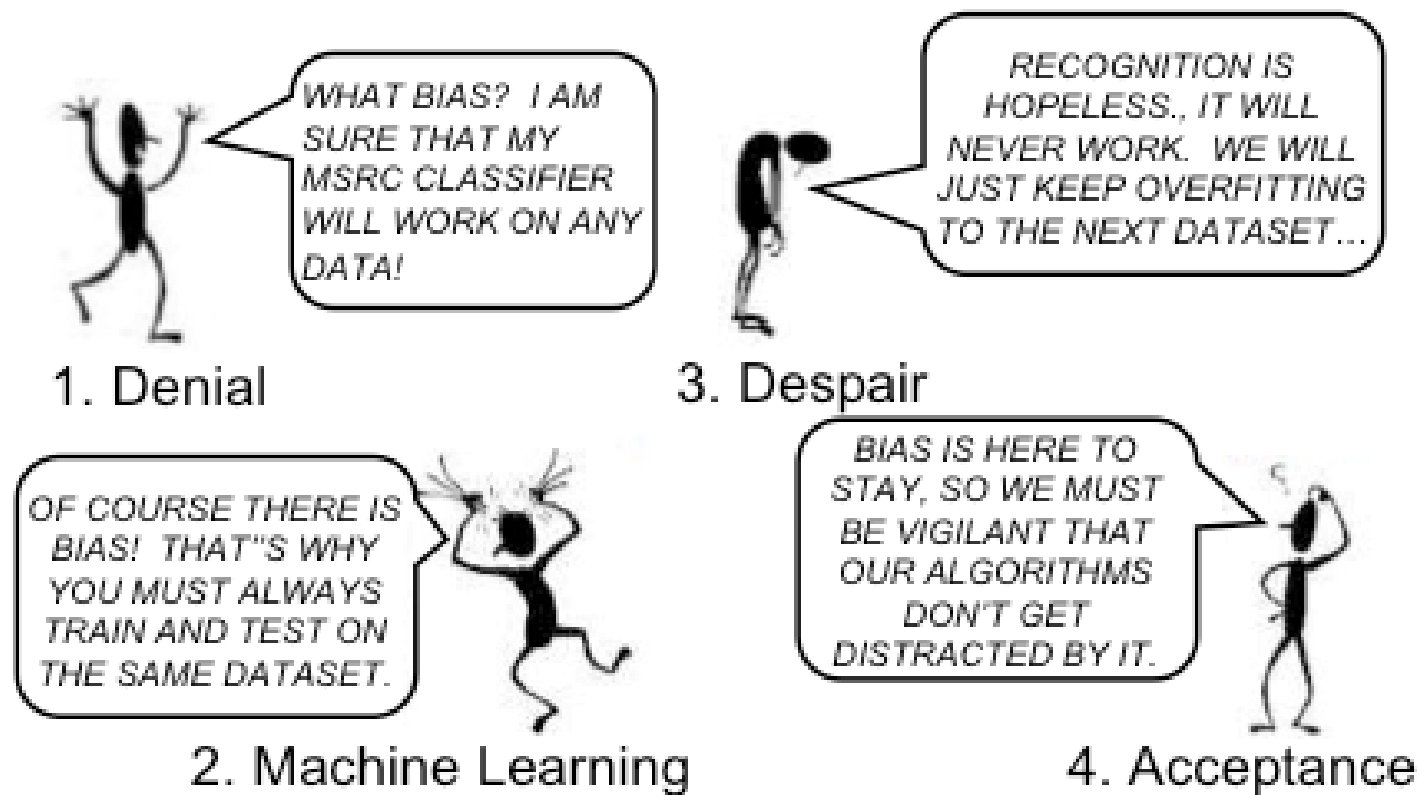


2. Machine Learning

Problem




Problem




Deep Learning-based Domain Adaptation

Deep Unsupervised domain adaptation by back propagation



- Computes $d = G_d(\mathbf{f}; \theta_d)$
- Is trained to predict **0** for **source** and **1** for **target**
- Therefore, the domain loss

is **low** for



is **higher** for





Fig. credit: Yaroslav Ganin
Yaroslav Ganin and Viktor Lempitsky
Unsupervised domain adaptation by back propagation
ICML 2015


Assumptions

- There are
- lots of labeled examples in source domain
- lots of unlabelled examples in target domain
- We want a deep neural network that does well on target domain


Deep Unsupervised domain adaptation



Deep Unsupervised domain adaptation



Deep Unsupervised domain adaptation



When trained on **source only**,
feature distributions **do not**
match.

$$S(\mathbf{f}) = \{G_f(\mathbf{x}; \theta_f) \mid \mathbf{x} \sim S(\mathbf{x})\}$$

$$T(\mathbf{f}) = \{G_f(\mathbf{x}; \theta_f) \mid \mathbf{x} \sim T(\mathbf{x})\}$$






Fig. credit: Yaroslav Ganin
Yaroslav Ganin and Viktor Lempitsky
Unsupervised domain adaptation by back propagation
ICML 2015

Deep Unsupervised domain adaptation



When trained on **source only**,
feature distributions **do not**
match.

Our goal is to get this:



Deep Unsupervised domain adaptation






Fig. credit: Yaroslav Ganin
Yaroslav Ganin and Viktor Lempitsky
Unsupervised domain adaptation by back propagation
ICML 2015

Deep Unsupervised domain adaptation




■ Computes $d = G_d(\mathbf{f}; \theta_d)$

Deep Unsupervised domain adaptation



- Computes $d = G_d(\mathbf{f}; \theta_d)$
- Is trained to predict **0** for **source** and **1** for **target**
- Therefore, the domain loss

is **low** for



is **higher** for



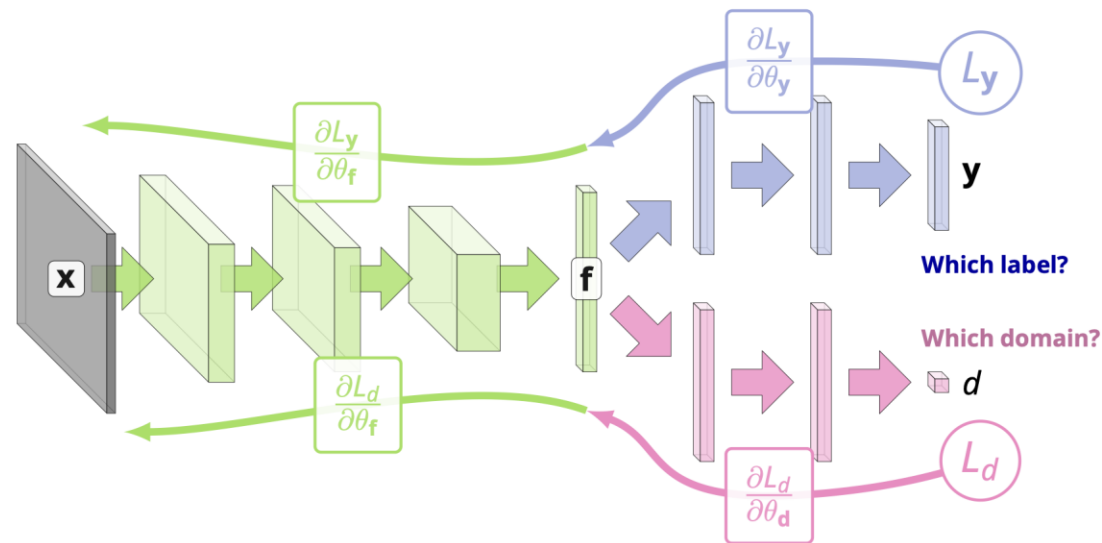


Fig. credit: Yaroslav Ganin
Yaroslav Ganin and Viktor Lempitsky
Unsupervised domain adaptation by back propagation
ICML 2015

Deep Unsupervised domain adaptation




Deep Unsupervised domain adaptation



Let's try *standard backpropagation*. **Emerging features** are:

- Discriminative (i.e. good for predicting y)
- Domain-discriminative (i.e. good for predicting d)


Deep Unsupervised domain adaptation



Let's now inject the **Gradient Reversal Layer**:

- Copies data without change at *fprop*
- Multiplies deltas by $-\lambda$ at *bprop*


Deep Unsupervised domain adaptation



Emerging features are now:


- Discriminative (i.e. good for predicting y)
- Domain-invariant (i.e. not good for predicting d)

Deep Unsupervised domain adaptation




- 1 Train **feature extractor** + **label predictor** on **source**
- 2 Train **feature extractor** + **domain classifier** on **source** + **target**
- 3 Use **feature extractor** + **label predictor** at test time

Deep Unsupervised domain adaptation



- 1 Train **feature extractor** + **label predictor** on **source**
- 2 Train **feature extractor** + **domain classifier** on **source** + **target**
- 3 Use **feature extractor** + **label predictor** at test time

Deep Unsupervised domain adaptation



Deep Unsupervised domain adaptation

METHOD	SOURCE	AMAZON	DSLR	WEBCAM
	TARGET	WEBCAM	WEBCAM	DSLR
GFK(PLS, PCA) (<i>GONG ET AL., 2013</i>)		.197	.497	.631
SA (<i>FERNANDO ET AL., 2013</i>)		.450	.648	.699
DLID (<i>S. CHOPRA & GOPALAN, 2013</i>)		.519	.782	.899
DDC (<i>TZENG ET AL., 2014</i>)		.618	.950	.985
DAN (<i>LONG & WANG, 2015</i>)		.685	.960	.990
SOURCE ONLY		.642	.961	.978
PROPOSED APPROACH		.730	.964	.992

Protocol: all of the methods above use

- all available **labeled source** samples
- all available **unlabeled target** samples

Domain Adaptation using Adaboost

Domain Adaptation Setting

- Example: We want to obtain sentiment analysis for movies
- We have a small amount of data for movies for the year 2020
- This data is insufficient to train for sentiment analysis
- We also have some additional training data for movies released in 1990-1995
- How can we best make use of the old training data

Sentiment Analysis
dataset for year 2020

Sentiment Analysis
dataset for years 1990-
1995

Domain Adaptation Setting

Sentiment Analysis
dataset for year
2020

Option 1: Train only with Y2020 data

Problem: Data not enough

Sentiment Analysis dataset for year
2020

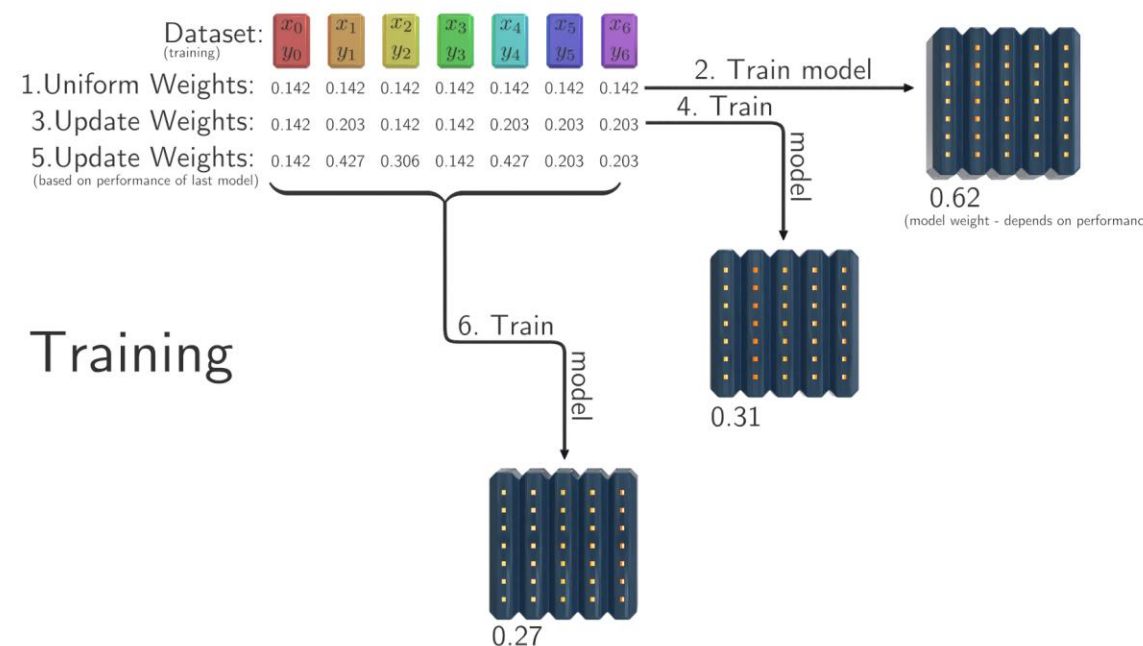
Sentiment Analysis dataset for years
1990-1995

Option 2: Put all data together and train

Problem: While some of the old data is relevant, some of it is not very relevant

How to Solve?

- Can we find out the data that is relevant from old data
- How do we go about it
- Consider that you have been taught Adaboost and you are a big fan of it :)
- Can you use Adaboost and solve the problem?



Solution: TrAdaboost

- Main idea:
- Let X_s be source dataset, (the old dataset)
- Let X_d be the new labeled target dataset (the new dataset)
- Train using all samples
- If a sample is misclassified, check which dataset it comes from
- If it is source dataset, then decrease its weight, if it is target dataset then increase its weight

Boosting for Transfer Learning

Wenyuan Dai DWYAK@APEX.SJTU.EDU.CN
Department of Computer Science and Engineering, Shanghai Jiao Tong University, China

Qiang Yang QYANG@CSE.UST.HK
Department of Computer Science, Hong Kong University of Science and Technology, Hong Kong

Gui-Rong Xue GRXUE@APEX.SJTU.EDU.CN
Yong Yu YYU@APEX.SJTU.EDU.CN
Department of Computer Science and Engineering, Shanghai Jiao Tong University, China

TrAdaBoost: ICML 2007

And it works

Table 4. The error rates when semi-supervised learning

Data Set	TSVM	TSVMt	TrAdaBoost(TSVM)
rec vs talk	0.059	0.040	0.021
rec vs sci	0.067	0.062	0.013
sci vs talk	0.173	0.106	0.075
auto vs aviation	0.043	0.103	0.038
real vs simulated	0.144	0.131	0.102
orgs vs people	0.358	0.292	0.248
orgs vs places	0.424	0.436	0.304
people vs places	0.307	0.225	0.179
edible vs poisonous	0.439	0.179	0.160

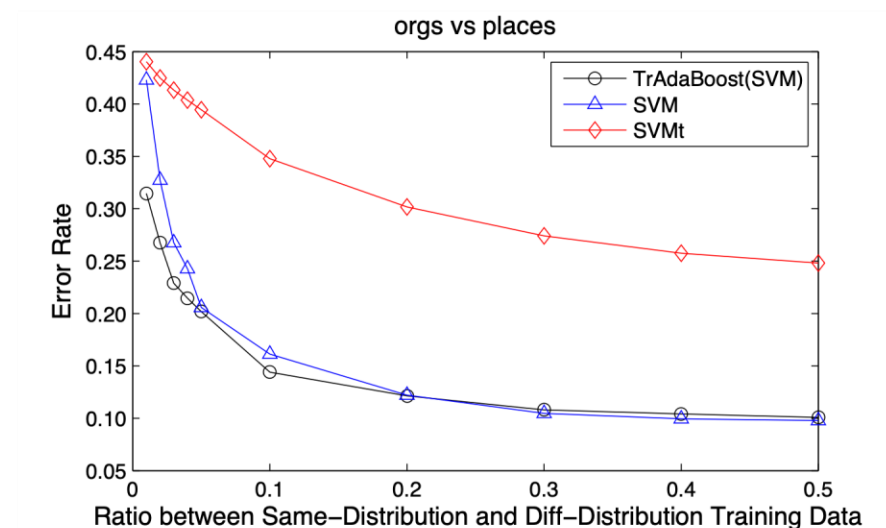


Figure 2. The error rate curves on **orgs vs places** data set for three classifiers **TrAdaBoost(SVM)**, **SVM** and **SVMt**

TrAdaboost Algorithm

TrAdaboost Algorithm

Input: The source dataset that is fully labeled (X_s, Y_s) consisting of n samples and the destination dataset X_d, Y_d consisting of m samples such that $m < n$, a baseline learner that provides a hypothesis $h(x) \rightarrow Y$. Initialise $\hat{w}_{0(n+1, \dots, n+m)} = [\frac{1}{m}, \dots, \frac{1}{m}]$. For $t = 1$ to N do,

TrAdaboost Algorithm

Input: The source dataset that is fully labeled (X_s, Y_s) consisting of n samples and the destination dataset X_d, Y_d consisting of m samples such that $m < n$, a baseline learner that provides a hypothesis $h(x) \rightarrow Y$. Initialise $\hat{w}_{0(n+1, \dots, n+m)} = [\frac{1}{m}, \dots, \frac{1}{m}]$. For $t = 1$ to N do,

1. Train model M_t with \hat{w}_t *weighted data*

TrAdaboost Algorithm

Input: The source dataset that is fully labeled (X_s, Y_s) consisting of n samples and the destination dataset X_d, Y_d consisting of m samples such that $m < n$, a baseline learner that provides a hypothesis $h(x) \rightarrow Y$. Initialise $\hat{w}_{0(n+1, \dots, n+m)} = [\frac{1}{m}, \dots, \frac{1}{m}]$. For $t = 1$ to N do,

1. Train model M_t with \hat{w}_t *weighted data*
2. Calculate weighted error over X_d

$$e_t = \frac{1}{\|\hat{w}_t\|_1} \sum_{i=n+1}^{n+m} w_{ti} L_{ti}$$

where L_{ti} is zero-one loss of exemplar i for model M_t

TrAdaboost Algorithm

Input: The source dataset that is fully labeled (X_s, Y_s) consisting of n samples and the destination dataset X_d, Y_d consisting of m samples such that $m < n$, a baseline learner that provides a hypothesis $h(x) \rightarrow Y$. Initialise $\hat{w}_{0(n+1, \dots, n+m)} = [\frac{1}{m}, \dots, \frac{1}{m}]$. For $t = 1$ to N do,

1. Train model M_t with \hat{w}_t *weighted data*
2. Calculate weighted error over X_d

$$e_t = \frac{1}{\|\hat{w}_t\|_1} \sum_{i=n+1}^{n+m} w_{ti} L_{ti}$$

where L_{ti} is zero-one loss of exemplar i for model M_t

3. Model weight

$$\alpha_t = \log \left(\frac{1 - e_t}{e_t} \right), \text{ and } \alpha = \frac{1}{(1 + \sqrt{2 \ln n / N})}$$

TrAdaboost Algorithm

Input: The source dataset that is fully labeled (X_s, Y_s) consisting of n samples and the destination dataset X_d, Y_d consisting of m samples such that $m < n$, a baseline learner that provides a hypothesis $h(x) \rightarrow Y$. Initialise $\hat{w}_{0(n+1, \dots, n+m)} = [\frac{1}{m}, \dots, \frac{1}{m}]$. For $t = 1$ to N do,

1. Train model M_t with \hat{w}_t *weighted data*
2. Calculate weighted error over X_d

$$e_t = \frac{1}{\|\hat{w}_t\|_1} \sum_{i=n+1}^{n+m} w_{ti} L_{ti}$$

where L_{ti} is zero-one loss of exemplar i for model M_t

3. Model weight

$$\alpha_t = \log \left(\frac{1 - e_t}{e_t} \right), \text{ and } \alpha = \frac{1}{(1 + \sqrt{2 \ln n / N})}$$

4. Update weights (only changes weights of exemplars it got wrong)

$$w_{t+1,i} = w_{ti} \exp(\alpha_t L_{ti}) \text{ for } i \in (n+1, n+m)$$

$$w_{t+1,i} = w_{ti} \exp(-\alpha L_{ti}) \text{ for } i \in (1, n)$$

Related Reading List

- Boosting for transfer learning, by Wenyuan Dai, Qiang Yang, Gui-Rong Xue, Yong Yu, ICML 2007
- Unsupervised domain adaptation by back propagation, Yaroslav Ganin and Viktor Lempitsky, ICML 2015